

IRIS Dataset

Machine Learning

Supervised Learning

Classification

- Logistic Regression / classification
- Decision Tree Classifier
- Random Forest Classifier
- K-N Neighbor Classifier
- Support Vector Machine - Linear, Ploly, RBF(Redial Basis Function)
- Naive Bayes - MultinomialNB, BurnoulliNB, GoussianNB

Regression

- Simple Linear Regression
- Multi- Linear Regression
- Decision Tree Regresor
- Random Forest Regressor
- K -N Neighbour Regressor
- Support Vector Machine Learning - Linear, polly, RBF - Redial Basis Function

Unsupervised Learning

- K-Mean Clustering
- Hierarchical Clustering

<https://www.javatpoint.com/regression-vs-classification-in-machine-learning#:~:text=The%20main%20difference%20between%20Regression%20and%20Classification>
<https://www.javatpoint.com/regression-vs-classification-in-machine-learning#:~:text=The%20main%20difference%20between%20Regression%20and%20Classification>



Import libraries

```
In [1]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

Read DataSet

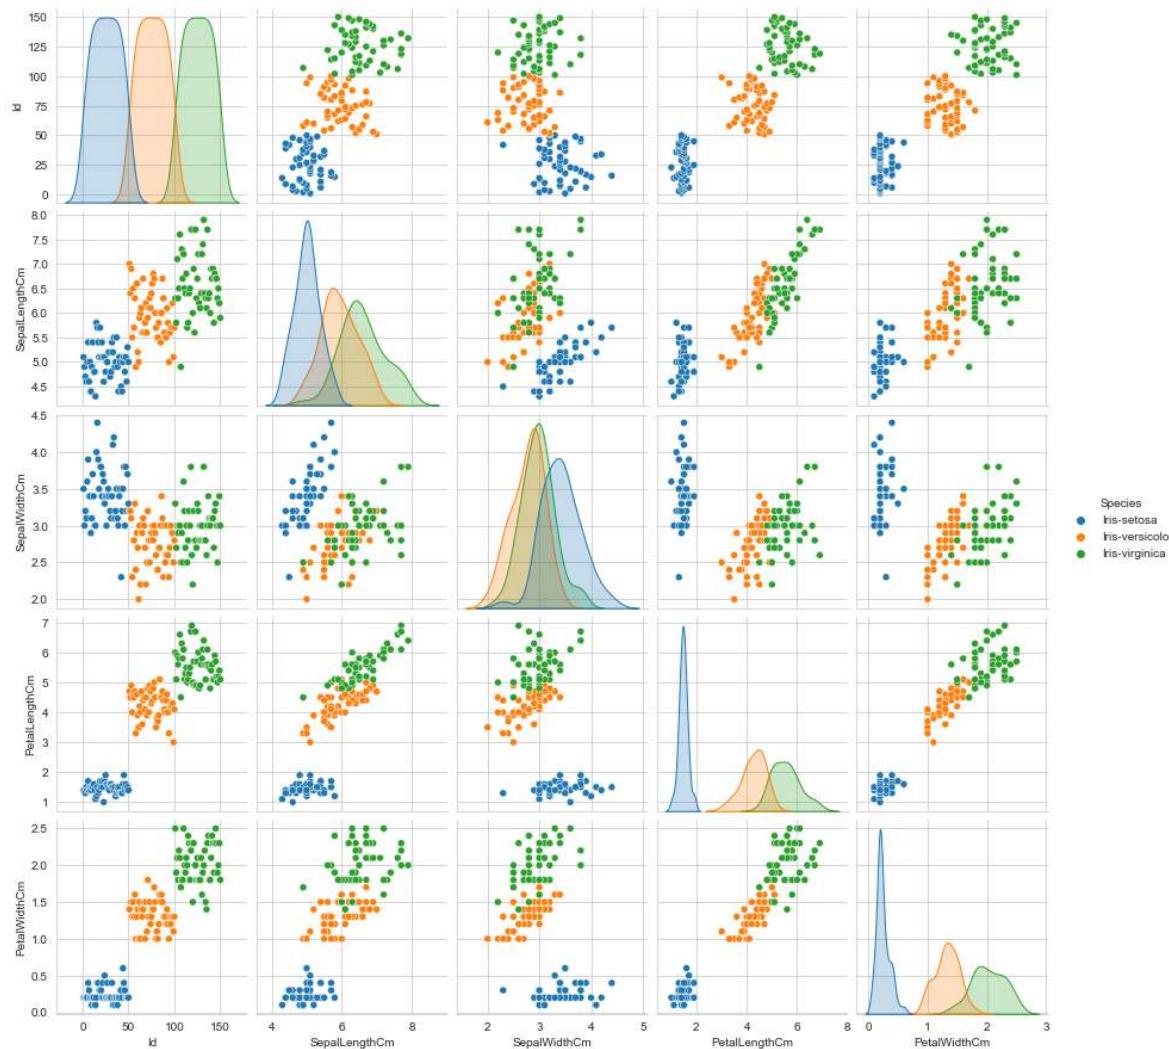
```
In [2]: df=pd.read_csv("C:\\Users\\omkan\\Desktop\\Iris1.csv")  
df1=df
```

```
In [3]: df.head(10)
```

Out[3]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa
3	4	4.6	3.1	1.5	0.2	Iris-setosa
4	5	5.0	3.6	1.4	0.2	Iris-setosa
5	6	5.4	3.9	1.7	0.4	Iris-setosa
6	7	4.6	3.4	1.4	0.3	Iris-setosa
7	8	5.0	3.4	1.5	0.2	Iris-setosa
8	9	4.4	2.9	1.4	0.2	Iris-setosa
9	10	4.9	3.1	1.5	0.1	Iris-setosa

```
In [4]: sns.set_style("whitegrid")
sns.pairplot(df,hue="Species");
plt.show()
```



Analyses

```
In [5]: df.corr()
```

Out[5]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
Id	1.000000	0.716676	-0.397729	0.882747	0.899759
SepalLengthCm	0.716676	1.000000	-0.109369	0.871754	0.817954
SepalWidthCm	-0.397729	-0.109369	1.000000	-0.420516	-0.356544
PetalLengthCm	0.882747	0.871754	-0.420516	1.000000	0.962757
PetalWidthCm	0.899759	0.817954	-0.356544	0.962757	1.000000

In [6]: df.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 6 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Id               150 non-null    int64  
 1   SepalLengthCm    150 non-null    float64 
 2   SepalWidthCm     150 non-null    float64 
 3   PetalLengthCm   150 non-null    float64 
 4   PetalWidthCm    150 non-null    float64 
 5   Species          150 non-null    object  
dtypes: float64(4), int64(1), object(1)
memory usage: 7.2+ KB
```

In [7]: df.describe()

Out[7]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	75.500000	5.843333	3.054000	3.758667	1.198667
std	43.445368	0.828066	0.433594	1.764420	0.763161
min	1.000000	4.300000	2.000000	1.000000	0.100000
25%	38.250000	5.100000	2.800000	1.600000	0.300000
50%	75.500000	5.800000	3.000000	4.350000	1.300000
75%	112.750000	6.400000	3.300000	5.100000	1.800000
max	150.000000	7.900000	4.400000	6.900000	2.500000

In [8]: df.isnull().sum()

Out[8]:

```
Id            0
SepalLengthCm 0
SepalWidthCm  0
PetalLengthCm 0
PetalWidthCm  0
Species        0
dtype: int64
```

In [9]: df.head(2)

Out[9]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa

```
In [10]: df.head(3)
```

Out[10]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	Iris-setosa
1	2	4.9	3.0	1.4	0.2	Iris-setosa
2	3	4.7	3.2	1.3	0.2	Iris-setosa

CLASSIFICATION

Logistic Regression

- In the case of Classification, we need to change Label/target/output into discrete form LabelEncoder

```
In [11]: df['Species'].value_counts()
```

```
Out[11]: Iris-setosa      50  
Iris-versicolor    50  
Iris-virginica     50  
Name: Species, dtype: int64
```

```
In [12]: from sklearn.preprocessing import LabelEncoder  
le=LabelEncoder()  
df['Species']=le.fit_transform(df['Species'])  
df['Species'].value_counts()
```

```
Out[12]: 0    50  
1    50  
2    50  
Name: Species, dtype: int64
```

```
In [13]: df.head(2)
```

Out[13]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	0
1	2	4.9	3.0	1.4	0.2	0

```
In [14]: df01=df[df['Species']!=2]
```

```
In [15]: df01['Species'].value_counts()
```

```
Out[15]: 0    50  
1    50  
Name: Species, dtype: int64
```

```
In [16]: df01.head()
```

```
Out[16]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	0
1	2	4.9	3.0	1.4	0.2	0
2	3	4.7	3.2	1.3	0.2	0
3	4	4.6	3.1	1.5	0.2	0
4	5	5.0	3.6	1.4	0.2	0

```
In [17]: df01.tail()
```

```
Out[17]:
```

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
95	96	5.7	3.0	4.2	1.2	1
96	97	5.7	2.9	4.2	1.3	1
97	98	6.2	2.9	4.3	1.3	1
98	99	5.1	2.5	3.0	1.1	1
99	100	5.7	2.8	4.1	1.3	1

```
In [18]: x=df01.iloc[:,1:5]      # independent variable  
y=df01['Species']           # Dependent variable  
x.shape,y.shape
```

```
Out[18]: ((100, 4), (100,))
```

```
In [19]: x.head(3)
```

```
Out[19]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
0	5.1	3.5	1.4	0.2
1	4.9	3.0	1.4	0.2
2	4.7	3.2	1.3	0.2

```
In [20]: y.head(3)
```

```
Out[20]: 0    0  
1    0  
2    0  
Name: Species, dtype: int32
```

```
In [21]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.25)  
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[21]: ((75, 4), (25, 4), (75,), (25,))
```

```
In [22]: x_train.head(2)
```

```
Out[22]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
18	5.7	3.8	1.7	0.3
95	5.7	3.0	4.2	1.2

```
In [23]: x_test.head(2)
```

```
Out[23]:
```

	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm
5	5.4	3.9	1.7	0.4
24	4.8	3.4	1.9	0.2

```
In [24]: y_train.head(2)
```

```
Out[24]: 18    0  
95    1  
Name: Species, dtype: int32
```

```
In [25]: y_test.head(2)
```

```
Out[25]: 5    0  
24   0  
Name: Species, dtype: int32
```

Import Model/algorith

```
In [26]: from sklearn.linear_model import LogisticRegression  
lr=LogisticRegression()
```

Train Model

```
In [27]: lr.fit(x_train,y_train)
```

```
Out[27]: LogisticRegression()
```

Predict Model

```
In [28]: y_pred_lr=lr.predict(x_test)  
y_pred_lr[:5],y_test.values[:5]
```

```
Out[28]: (array([0, 0, 1, 1, 0]), array([0, 0, 1, 1, 0]))
```

Evaluation

```
In [29]: print(lr.score(x_train,y_train))  
print(lr.score(x_test,y_test))
```

```
1.0
```

```
1.0
```

```
In [30]: from sklearn.metrics import confusion_matrix,classification_report,accuracy_score  
print(confusion_matrix(y_pred_lr,y_test))  
print(classification_report(y_pred_lr,y_test))  
print(f'model_score- {lr.score(x_test,y_test)} ')  
print(f'accuracy_score- {accuracy_score(y_pred_lr,y_test)} ')
```

```
[[ 9  0]  
 [ 0 16]]
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16

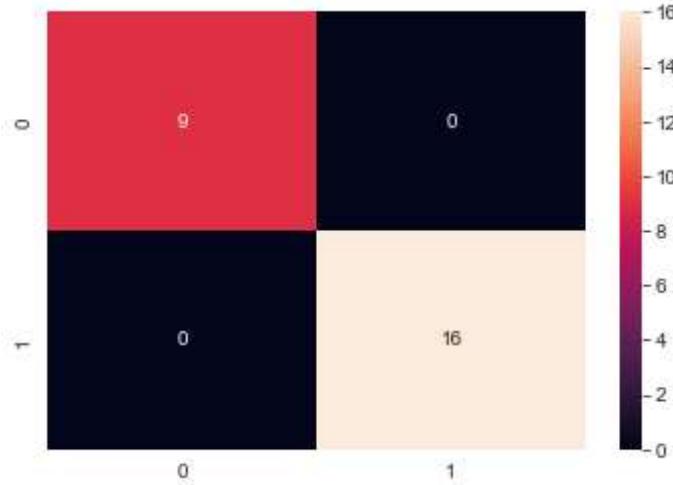
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

```
model_score- 1.0
```

```
accuracy_score- 1.0
```

Vsualization

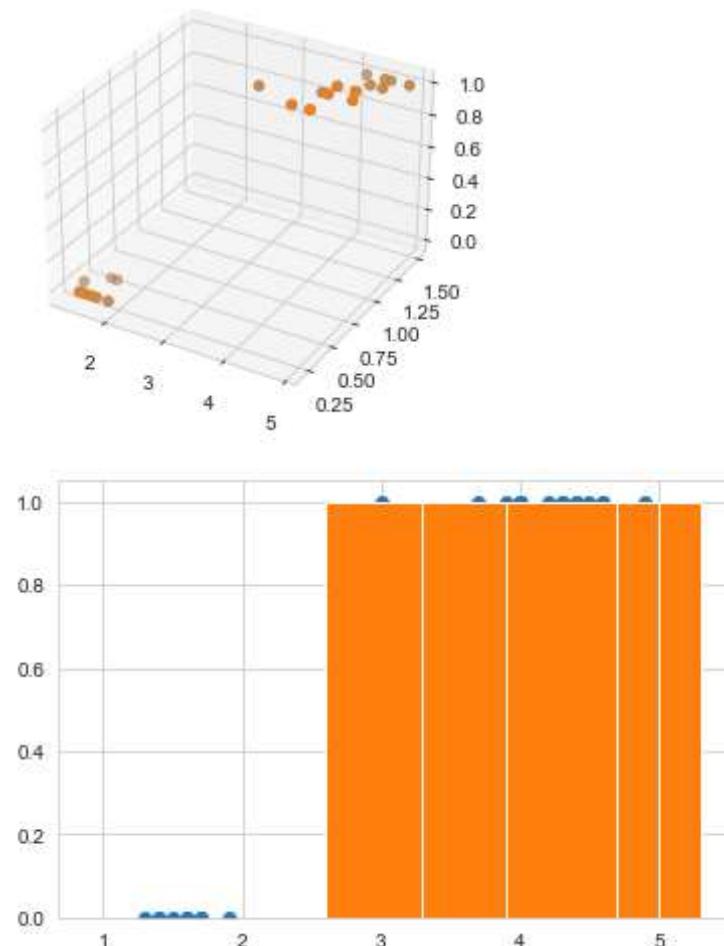
```
In [31]: cm=confusion_matrix(y_pred_lr,y_test)
sns.heatmap(cm,annot=True)
plt.show()
```



<https://datascience.stackexchange.com/questions/65839/macro-average-and-weighted-average-meaning-in-classification-report>
(<https://datascience.stackexchange.com/questions/65839/macro-average-and-weighted-average-meaning-in-classification-report>)

<https://www.kaggle.com/getting-started/175898> (<https://www.kaggle.com/getting-started/175898>)

```
In [32]: ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_lr,'black')
plt.show()
plt.scatter(x_test['PetalLengthCm'],y_test)
plt.bar(x_test['PetalLengthCm'],y_pred_lr)
plt.show()
```



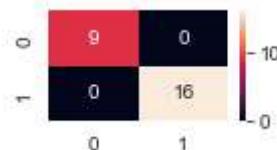
Decision Tree Classifier

```
In [33]: from sklearn.tree import DecisionTreeClassifier
dtc=DecisionTreeClassifier()
dtc.fit(x_train,y_train)
y_pred_dtc=dtc.predict(x_test)
print(f'Predicted_y{y_pred_dtc[:5]} Actual_y{y_test.values[:5]}')
print(confusion_matrix(y_pred_dtc,y_test))
from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
cm=confusion_matrix(y_pred_dtc,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_dtc,y_test))
print(f'model_score-{dtc.score(x_test,y_test)}')
print(f'accuracy_score-{accuracy_score(y_pred_dtc,y_test)}')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_dtc,'black')
plt.show()
```

Predicted_y[0 0 1 1 0] Actual_y[0 0 1 1 0]

[[9 0]
[0 16]]

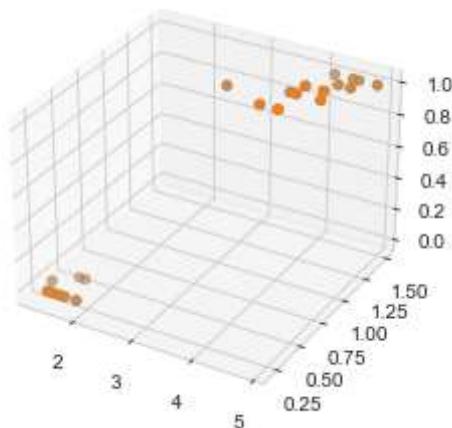
AxesSubplot(0.125,0.125;0.62x0.755)



	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

model_score- 1.0

accuracy_score- 1.0



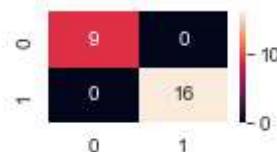
RandomForestClassifier

<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>
[\(https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html)

```
In [34]: from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
y_pred_rfc=rfc.predict(x_test)
print(f' predicted_y-{y_pred_rfc} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_rfc,y_test))
cm=confusion_matrix(y_pred_rfc,y_test)
plt.figure(figsize=(2,1))
sns.heatmap(cm,annot=True)
plt.show()
print(classification_report(y_pred_rfc,y_test))
print(f'model_score- {rfc.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_rfc,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_rfc,'black')
plt.show()
```

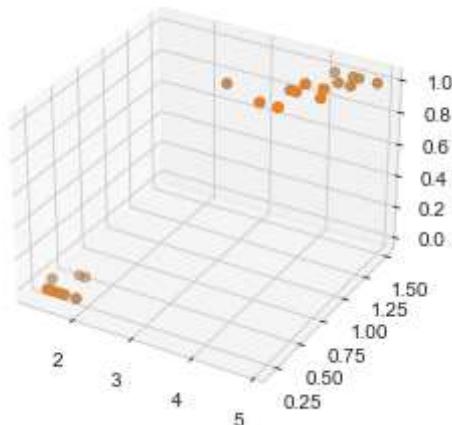
predicted_y-[0 0 1 1 0 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0] actual_y-[0 0 1 1 0 1 1 1 1 0 1 0 1 1 0 1 1 1 1 0 1 0]

[[9 0]
[0 16]]



	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

model_score- 1.0
accuracy_score- 1.0



K-Nearest Neighbours

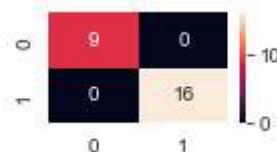
<https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html>
[\(https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html\)](https://scikit-learn.org/stable/modules/generated/sklearn.neighbors.KNeighborsClassifier.html)

```
In [35]: from sklearn.neighbors import KNeighborsClassifier
knc=KNeighborsClassifier()
knc.fit(x_train,y_train)
y_pred_knc=knc.predict(x_test)
print(f'Predicted_y{y_pred_knc[:5]} Actual_y{y_test.values[:5]}')
print(confusion_matrix(y_pred_knc,y_test))
from sklearn.metrics import confusion_matrix,classification_report,accuracy_score
cm=confusion_matrix(y_pred_knc,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_knc,y_test))
print(f'model_score-{dtc.score(x_test,y_test)}')
print(f'accuracy_score-{accuracy_score(y_pred_knc,y_test)}')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_knc,'black')
plt.show()
```

Predicted_y[0 0 1 1 0] Actual_y[0 0 1 1 0]

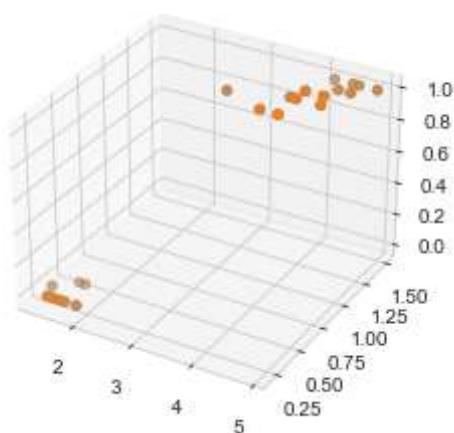
[[9 0]
[0 16]]

AxesSubplot(0.125,0.125;0.62x0.755)



	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

model_score- 1.0
accuracy_score- 1.0

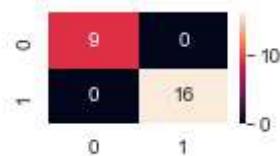


Naive Bayes

- gaussian
- bernoulli
- multinomial

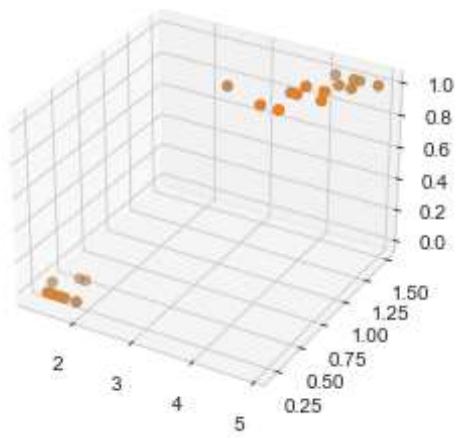
```
In [36]: print( 'Naive Bayes - gaussian')
from sklearn.naive_bayes import GaussianNB
gnb=GaussianNB()
gnb.fit(x_train,y_train)
y_pred_gnb=gnb.predict(x_test)
print(f' predicted_y-{y_pred_gnb} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_gnb,y_test))
cm=confusion_matrix(y_pred_gnb,y_test)
plt.figure(figsize=(2,1))
sns.heatmap(cm,annot=True)
plt.show()
print(classification_report(y_pred_gnb,y_test))
print(f'model_score- {gnb.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_gnb,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_gnb,'black')
plt.show()
```

Naive Bayes - gaussian
predicted_y-[0 0 1 1 0 1 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
actual_y-[0
0 1 1 0 1 1 1 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
[[9 0]
[0 16]]



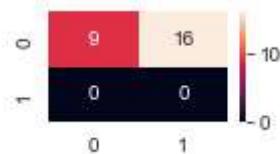
	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

model_score- 1.0
accuracy_score- 1.0



```
In [37]: print( 'Naive Bayes - Bernoulli')
from sklearn.naive_bayes import BernoulliNB
bnb=BernoulliNB()
bnb.fit(x_train,y_train)
y_pred_bnb=bnb.predict(x_test)
print(f' predicted_y-{y_pred_bnb} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_bnb,y_test))
cm=confusion_matrix(y_pred_bnb,y_test)
plt.figure(figsize=(2,1))
sns.heatmap(cm,annot=True)
plt.show()
print(classification_report(y_pred_bnb,y_test))
print(f'model_score- {bnb.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_bnb,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_bnb,'black')
plt.show()
plt.scatter(x_test['PetalLengthCm'],y_test)
plt.scatter(x_test['PetalLengthCm'],y_pred_bnb)
```

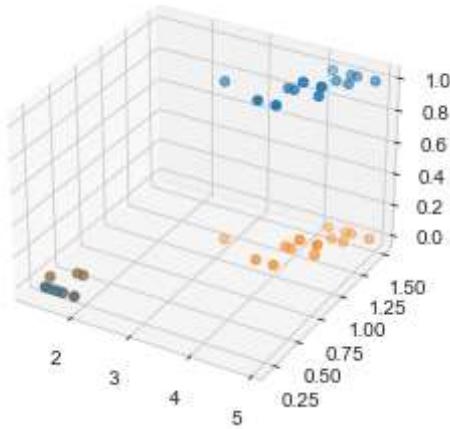
Naive Bayes - Bernoulli
predicted_y-[0 0] actual_y-[0
0 1 1 0 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
[[9 16]
[0 0]]



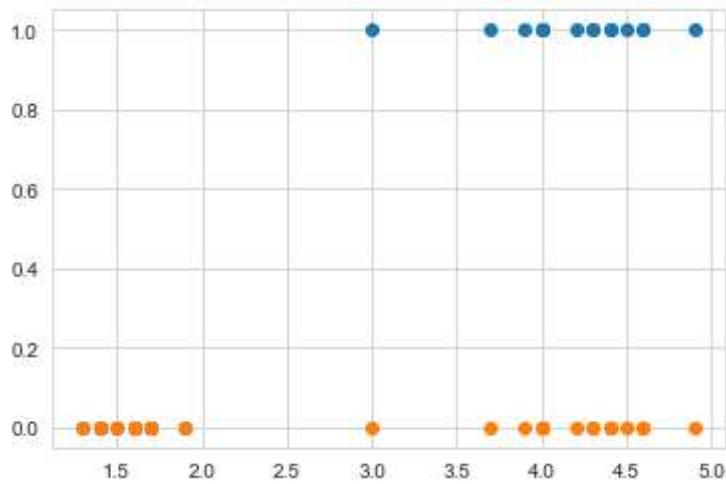
	precision	recall	f1-score	support
0	1.00	0.36	0.53	25
1	0.00	0.00	0.00	0
accuracy			0.36	25
macro avg	0.50	0.18	0.26	25
weighted avg	1.00	0.36	0.53	25

model_score- 0.36
accuracy_score- 0.36

```
C:\Users\omkan\anaconda3\lib\site-packages\sklearn\metrics\_classification.p
y:1248: UndefinedMetricWarning: Recall and F-score are ill-defined and being
set to 0.0 in labels with no true samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\omkan\anaconda3\lib\site-packages\sklearn\metrics\_classification.p
y:1248: UndefinedMetricWarning: Recall and F-score are ill-defined and being
set to 0.0 in labels with no true samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
C:\Users\omkan\anaconda3\lib\site-packages\sklearn\metrics\_classification.p
y:1248: UndefinedMetricWarning: Recall and F-score are ill-defined and being
set to 0.0 in labels with no true samples. Use `zero_division` parameter to
control this behavior.
    _warn_prf(average, modifier, msg_start, len(result))
```

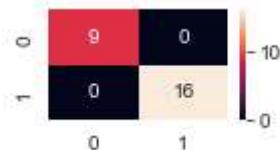


Out[37]: <matplotlib.collections.PathCollection at 0x1eb9cc6a5b0>



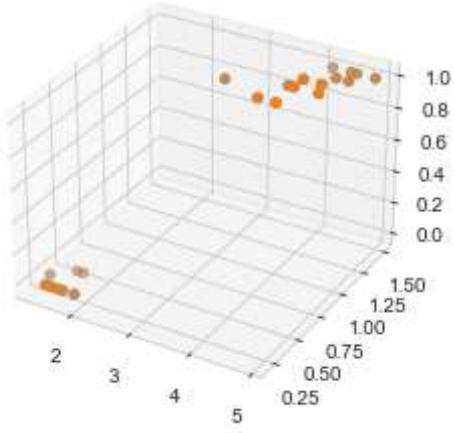
```
In [38]: print( 'Naive Bayes - Multinomial')
from sklearn.naive_bayes import MultinomialNB
mnb=MultinomialNB()
mnb.fit(x_train,y_train)
y_pred_mnb=mnb.predict(x_test)
print(f' predicted_y-{y_pred_mnb} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_mnb,y_test))
cm=confusion_matrix(y_pred_mnb,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_mnb,y_test))
print(f'model_score- {mnb.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_mnb,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_mnb,'black')
plt.show()
```

Naive Bayes - Multinomial
predicted_y-[0 0 1 1 0 1 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
actual_y-[0
0 1 1 0 1 1 1 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
[[9 0]
[0 16]]
AxesSubplot(0.125,0.125;0.62x0.755)



	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

model_score- 1.0
accuracy_score- 1.0



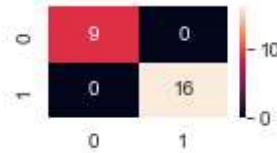
Support Vector Machine - Classification

- kernel - poly, linear, rbf (Radial Basis Function)

```
In [39]: from sklearn.svm import SVC
```

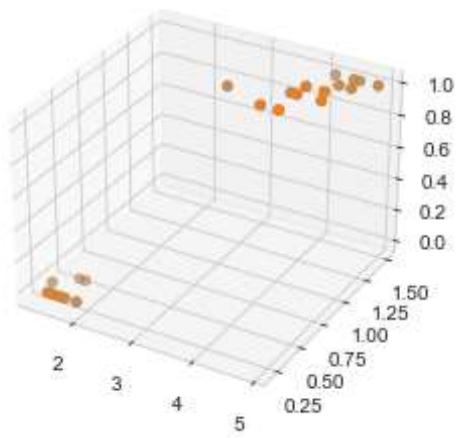
```
In [40]: print(' svc - poly')
svcp=SVC(kernel='poly', degree=2, C=.1)
svcp.fit(x_train,y_train)
y_pred_svcp=svcp.predict(x_test)
print(f' predicted_y-{y_pred_svcp} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_svcp,y_test))
cm=confusion_matrix(y_pred_svcp,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_svcp,y_test))
print(f'model_score- {svcp.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_svcp,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_svcp,'black')
plt.show()
```

svc - poly
 predicted_y-[0 0 1 1 0 1 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
 actual_y-[0
 0 1 1 0 1 1 1 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
 [[9 0]
 [0 16]]
 AxesSubplot(0.125,0.125;0.62x0.755)



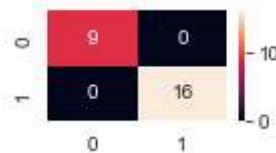
	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

model_score- 1.0
 accuracy_score- 1.0



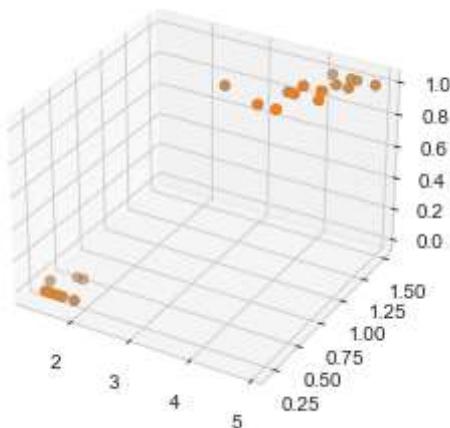
```
In [41]: print(' svc - linear')
svcl=SVC(kernel='linear', degree=2, C=.1).fit(x_train,y_train)
y_pred_svcl=svcl.predict(x_test)
print(f' predicted_y-{y_pred_svcl} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_svcl,y_test))
cm=confusion_matrix(y_pred_svcl,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_svcl,y_test))
print(f'model_score- {svcl.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_svcl,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_svcl,'black')
plt.show()
```

```
svc - linear
predicted_y-[ 0  0  1  1  0  1  1  1  1  0  1  0  1  0  1  1  0  1  1  1  1  0  1  0] actual_y-[ 0
0  1  1  0  1  1  1  1  0  1  0  1  1  0  1  1  1  1  0  1  0]
[[ 9  0]
 [ 0 16]]
AxesSubplot(0.125,0.125;0.62x0.755)
```



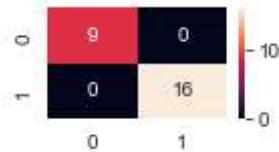
	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

```
model_score- 1.0
accuracy_score- 1.0
```



```
In [42]: print(' svc - rbf')
svcrbf=SVC(kernel='rbf', degree=2, C=.1)
svcrbf.fit(x_train,y_train)
y_pred_rbf=svcl.predict(x_test)
print(f' predicted_y-{y_pred_rbf} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_rbf,y_test))
cm=confusion_matrix(y_pred_rbf,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_rbf,y_test))
print(f'model_score- {svcrbf.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_rbf,y_test)} ')
```

```
svc - rbf
predicted_y-[0 0 1 1 0 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0] actual_y-[0
0 1 1 0 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
[[ 9  0]
 [ 0 16]]
AxesSubplot(0.125,0.125;0.62x0.755)
```



	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

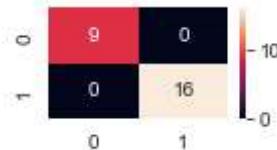
```
model_score- 1.0
accuracy_score- 1.0
```

```
In [43]: print('linearSVC')
from sklearn.svm import SVC
lsvc=SVC(C=.1, degree=.2)
```

```
linearSVC
```

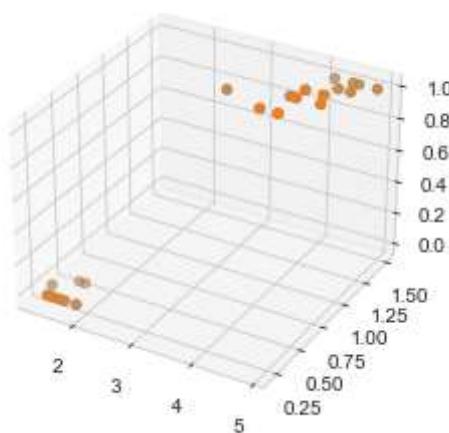
```
In [44]: lssvc.fit(x_train,y_train)
y_pred_lsvc=lssvc.predict(x_test)
print(f' predicted_y-{y_pred_lsvc} actual_y-{y_test.values}')
print(confusion_matrix(y_pred_lsvc,y_test))
cm=confusion_matrix(y_pred_lsvc,y_test)
plt.figure(figsize=(2,1))
print(sns.heatmap(cm,annot=True))
plt.show()
print(classification_report(y_pred_lsvc,y_test))
print(f'model_score- {lssvc.score(x_test,y_test)} ')
print(f'accuracy_score- {accuracy_score(y_pred_lsvc,y_test)} ')
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test[ 'PetalLengthCm'],x_test[ 'PetalWidthCm'],y_test)
ax.scatter3D(x_test[ 'PetalLengthCm'],x_test[ 'PetalWidthCm'],y_pred_lsvc,'black')
plt.show()
```

```
predicted_y-[0 0 1 1 0 1 1 1 1 0 1 0 1 0 1 1 0 1 1 1 0 1 0]
actual_y-[0
0 1 1 0 1 1 1 1 0 1 0 1 1 0 1 1 1 1 0 1 0]
[[ 9  0]
 [ 0 16]]
AxesSubplot(0.125,0.125;0.62x0.755)
```



	precision	recall	f1-score	support
0	1.00	1.00	1.00	9
1	1.00	1.00	1.00	16
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

```
model_score- 1.0
accuracy_score- 1.0
```



<https://datascience.stackexchange.com/questions/12355/support-vector-classification-kernels-linear-poly-rbf-has-all-same-score>
[\(https://datascience.stackexchange.com/questions/12355/support-vector-classification-kernels-linear-poly-rbf-has-all-same-score\)](https://datascience.stackexchange.com/questions/12355/support-vector-classification-kernels-linear-poly-rbf-has-all-same-score)

Neural Networks - classification

In [45]: `import tensorflow as tf
from tensorflow.keras import Sequential`

In [46]: `x_train.shape,x_test.shape,y_train.shape,y_test.shape`

Out[46]: `((75, 4), (25, 4), (75,), (25,))`

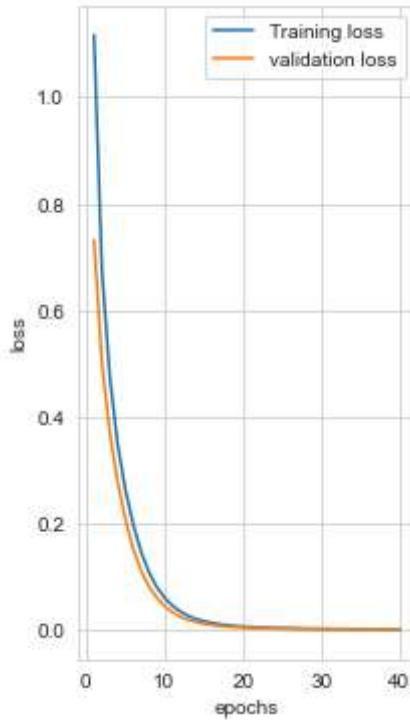
```
In [47]: inputs = tf.keras.Input(shape=(4,))
x=tf.keras.layers.Dense(256,activation='relu')(inputs)
x=tf.keras.layers.Dense(256,activation='relu')(x)
# two hidden input layers and 256 node
outputs=tf.keras.layers.Dense(3,activation = 'softmax')(x)
# one hidden output Layers and 3 node
model = tf.keras.Model(inputs,outputs)
model.compile(optimizer = 'adam',loss='sparse_categorical_crossentropy',metrics=['accuracy'])
batch_size=32
epochs = 40 # batch size * number of iterations = epoch
history = model.fit(x_train,y_train,validation_split=0.2,batch_size=batch_size,epochs=epochs)
#model.fit(x_train,y_train,validation_split=0.2,batch_size=batch_size,epochs=epochs)
```

Epoch 1/40
2/2 [=====] - 1s 153ms/step - loss: 1.1161 - accuracy: 0.2333 - val_loss: 0.7317 - val_accuracy: 0.4667
Epoch 2/40
2/2 [=====] - 0s 19ms/step - loss: 0.6763 - accuracy: 0.6833 - val_loss: 0.5016 - val_accuracy: 1.0000
Epoch 3/40
2/2 [=====] - 0s 16ms/step - loss: 0.4771 - accuracy: 1.0000 - val_loss: 0.3722 - val_accuracy: 1.0000
Epoch 4/40
2/2 [=====] - 0s 26ms/step - loss: 0.3496 - accuracy: 1.0000 - val_loss: 0.2791 - val_accuracy: 1.0000
Epoch 5/40
2/2 [=====] - 0s 24ms/step - loss: 0.2631 - accuracy: 1.0000 - val_loss: 0.2077 - val_accuracy: 1.0000
Epoch 6/40
2/2 [=====] - 0s 25ms/step - loss: 0.1985 - accuracy: 1.0000 - val_loss: 0.1527 - val_accuracy: 1.0000
Epoch 7/40
2/2 [=====] - 0s 24ms/step - loss: 0.1475 - accuracy: 1.0000 - val_loss: 0.1118 - val_accuracy: 1.0000
Epoch 8/40
2/2 [=====] - 0s 27ms/step - loss: 0.1089 - accuracy: 1.0000 - val_loss: 0.0825 - val_accuracy: 1.0000
Epoch 9/40
2/2 [=====] - 0s 32ms/step - loss: 0.0814 - accuracy: 1.0000 - val_loss: 0.0612 - val_accuracy: 1.0000
Epoch 10/40
2/2 [=====] - 0s 24ms/step - loss: 0.0610 - accuracy: 1.0000 - val_loss: 0.0456 - val_accuracy: 1.0000
Epoch 11/40
2/2 [=====] - 0s 24ms/step - loss: 0.0463 - accuracy: 1.0000 - val_loss: 0.0338 - val_accuracy: 1.0000
Epoch 12/40
2/2 [=====] - 0s 27ms/step - loss: 0.0351 - accuracy: 1.0000 - val_loss: 0.0254 - val_accuracy: 1.0000
Epoch 13/40
2/2 [=====] - 0s 28ms/step - loss: 0.0271 - accuracy: 1.0000 - val_loss: 0.0195 - val_accuracy: 1.0000
Epoch 14/40
2/2 [=====] - 0s 25ms/step - loss: 0.0213 - accuracy: 1.0000 - val_loss: 0.0153 - val_accuracy: 1.0000
Epoch 15/40
2/2 [=====] - 0s 25ms/step - loss: 0.0171 - accuracy: 1.0000 - val_loss: 0.0123 - val_accuracy: 1.0000
Epoch 16/40
2/2 [=====] - 0s 29ms/step - loss: 0.0140 - accuracy: 1.0000 - val_loss: 0.0100 - val_accuracy: 1.0000
Epoch 17/40
2/2 [=====] - 0s 24ms/step - loss: 0.0115 - accuracy: 1.0000 - val_loss: 0.0082 - val_accuracy: 1.0000
Epoch 18/40
2/2 [=====] - 0s 19ms/step - loss: 0.0096 - accuracy: 1.0000 - val_loss: 0.0069 - val_accuracy: 1.0000
Epoch 19/40
2/2 [=====] - 0s 29ms/step - loss: 0.0083 - accuracy: 1.0000 - val_loss: 0.0059 - val_accuracy: 1.0000

```
Epoch 20/40
2/2 [=====] - 0s 24ms/step - loss: 0.0071 - accuracy: 1.0000 - val_loss: 0.0051 - val_accuracy: 1.0000
Epoch 21/40
2/2 [=====] - 0s 20ms/step - loss: 0.0062 - accuracy: 1.0000 - val_loss: 0.0044 - val_accuracy: 1.0000
Epoch 22/40
2/2 [=====] - 0s 26ms/step - loss: 0.0055 - accuracy: 1.0000 - val_loss: 0.0039 - val_accuracy: 1.0000
Epoch 23/40
2/2 [=====] - 0s 16ms/step - loss: 0.0050 - accuracy: 1.0000 - val_loss: 0.0035 - val_accuracy: 1.0000
Epoch 24/40
2/2 [=====] - 0s 24ms/step - loss: 0.0045 - accuracy: 1.0000 - val_loss: 0.0032 - val_accuracy: 1.0000
Epoch 25/40
2/2 [=====] - 0s 24ms/step - loss: 0.0041 - accuracy: 1.0000 - val_loss: 0.0029 - val_accuracy: 1.0000
Epoch 26/40
2/2 [=====] - 0s 32ms/step - loss: 0.0038 - accuracy: 1.0000 - val_loss: 0.0027 - val_accuracy: 1.0000
Epoch 27/40
2/2 [=====] - 0s 33ms/step - loss: 0.0035 - accuracy: 1.0000 - val_loss: 0.0025 - val_accuracy: 1.0000
Epoch 28/40
2/2 [=====] - 0s 25ms/step - loss: 0.0032 - accuracy: 1.0000 - val_loss: 0.0023 - val_accuracy: 1.0000
Epoch 29/40
2/2 [=====] - 0s 24ms/step - loss: 0.0031 - accuracy: 1.0000 - val_loss: 0.0021 - val_accuracy: 1.0000
Epoch 30/40
2/2 [=====] - 0s 24ms/step - loss: 0.0029 - accuracy: 1.0000 - val_loss: 0.0020 - val_accuracy: 1.0000
Epoch 31/40
2/2 [=====] - 0s 24ms/step - loss: 0.0027 - accuracy: 1.0000 - val_loss: 0.0019 - val_accuracy: 1.0000
Epoch 32/40
2/2 [=====] - 0s 24ms/step - loss: 0.0026 - accuracy: 1.0000 - val_loss: 0.0018 - val_accuracy: 1.0000
Epoch 33/40
2/2 [=====] - 0s 24ms/step - loss: 0.0024 - accuracy: 1.0000 - val_loss: 0.0017 - val_accuracy: 1.0000
Epoch 34/40
2/2 [=====] - 0s 24ms/step - loss: 0.0023 - accuracy: 1.0000 - val_loss: 0.0016 - val_accuracy: 1.0000
Epoch 35/40
2/2 [=====] - 0s 22ms/step - loss: 0.0022 - accuracy: 1.0000 - val_loss: 0.0016 - val_accuracy: 1.0000
Epoch 36/40
2/2 [=====] - 0s 24ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0015 - val_accuracy: 1.0000
Epoch 37/40
2/2 [=====] - 0s 32ms/step - loss: 0.0021 - accuracy: 1.0000 - val_loss: 0.0014 - val_accuracy: 1.0000
Epoch 38/40
2/2 [=====] - 0s 24ms/step - loss: 0.0020 - accuracy: 1.0000 - val_loss: 0.0014 - val_accuracy: 1.0000
```

```
Epoch 39/40
2/2 [=====] - 0s 24ms/step - loss: 0.0019 - accurac
y: 1.0000 - val_loss: 0.0013 - val_accuracy: 1.0000
Epoch 40/40
2/2 [=====] - 0s 24ms/step - loss: 0.0018 - accurac
y: 1.0000 - val_loss: 0.0013 - val_accuracy: 1.0000
```

```
In [48]: plt.figure(figsize = (3,6))
epoch_range = range(1,epochs +1)
train_loss = history.history['loss']
val_loss = history.history['val_loss']
plt.plot(epoch_range,train_loss,label='Training loss')
plt.plot(epoch_range,val_loss, label = 'validation loss')
plt.xlabel('epochs')
plt.ylabel('loss')
plt.legend()
plt.show()
```



```
In [49]: loss,acc=model.evaluate(x_test,y_test)
print('loss',loss)
print('acc',acc)
```

```
1/1 [=====] - 0s 16ms/step - loss: 0.0029 - accurac
y: 1.0000
loss 0.0029290372040122747
acc 1.0
```

REGRESSION

Independent variable and dependent variable In Regression

```
In [50]: x=df.iloc[:,3:4]      # independent variable  
y=df.iloc[:,4:5]      # dependent variable
```

```
In [51]: x.head()
```

Out[51]:

	PetalLengthCm
0	1.4
1	1.4
2	1.3
3	1.5
4	1.4

```
In [52]: y.head()
```

Out[52]:

	PetalWidthCm
0	0.2
1	0.2
2	0.2
3	0.2
4	0.2

```
In [53]: x.shape, y.shape
```

Out[53]: ((150, 1), (150, 1))

split the data set

```
In [54]: from sklearn.model_selection import train_test_split
```

```
In [55]: x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=.8)
```

```
In [56]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

Out[56]: ((120, 1), (30, 1), (120, 1), (30, 1))

```
In [57]: x_train.head(2)
```

Out[57]:

	PetalLengthCm
91	4.6
77	5.0

```
In [58]: x_test.head(2)
```

Out[58]:

	PetalLengthCm
98	3.0
37	1.5

```
In [59]: y_train.head(2)
```

Out[59]:

	PetalWidthCm
91	1.4
77	1.7

```
In [60]: y_test.head(2)
```

Out[60]:

	PetalWidthCm
98	1.1
37	0.1

In []:

SIMPLE LINEAR REGRESSION

Import the model/algorithm

```
In [61]: from sklearn.linear_model import LinearRegression # as slr  
slr=LinearRegression()
```

Train the model

```
In [62]: slr.fit(x_train,y_train)
```

```
Out[62]: LinearRegression()
```

Predict

```
In [63]: y_pred=slr.predict(x_test)
```

```
In [64]: x_test.values[:5]
```

```
Out[64]: array([[3. ],
   [1.5],
   [5.8],
   [5.4],
   [5.2]])
```

```
In [65]: y_pred[:5]
```

```
Out[65]: array([[0.87412307],
   [0.27081288],
   [2.00030209],
   [1.83941937],
   [1.75897802]])
```

```
In [66]: y_test.values[:5]
```

```
Out[66]: array([[1.1],
   [0.1],
   [2.2],
   [2.1],
   [2.3]])
```

evaluation

```
In [67]: print(slr.score(x_train,y_train))
print(slr.score(x_test,y_test))
```

```
0.9297411171423415
0.9131168764240425
```

```
In [68]: from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
print(mean_absolute_error(y_test,y_pred))
print(mean_squared_error(y_test,y_pred))
print(r2_score(y_test,y_pred))
```

```
0.19567010999121745
0.0703029274935456
0.9131168764240425
```

visualization

```
In [69]: import matplotlib.pyplot as plt
```

```
In [70]: slr.predict([[3]])
```

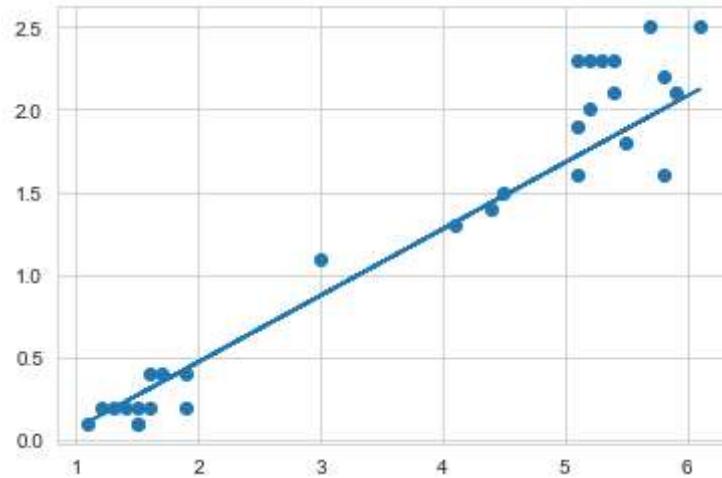
```
Out[70]: array([[0.87412307]])
```

```
In [71]: a=slr.coef_
b=slr.intercept_
```

```
In [72]: y =a*3+b    #  $y=ax+b$ 
y
```

```
Out[72]: array([[0.87412307]])
```

```
In [73]: plt.scatter(x_test,y_test)
plt.plot(x_test,y_pred)
plt.show()
```



MULTILINEAR REGRESSION

```
In [74]: df.head()
```

Out[74]:

	Id	SepalLengthCm	SepalWidthCm	PetalLengthCm	PetalWidthCm	Species
0	1	5.1	3.5	1.4	0.2	0
1	2	4.9	3.0	1.4	0.2	0
2	3	4.7	3.2	1.3	0.2	0
3	4	4.6	3.1	1.5	0.2	0
4	5	5.0	3.6	1.4	0.2	0

Divide dataset into independent and dependent variable

```
In [75]: x=df[['PetalLengthCm','PetalWidthCm']] # Independent variable  
y=df['SepalLengthCm'] # Dependent variable  
x.shape,y.shape
```

Out[75]: ((150, 2), (150,))

Split independent & dependent variables into train dataset and test dataset

```
In [76]: from sklearn.model_selection import train_test_split  
  
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size=.8)  
  
x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

Out[76]: ((120, 2), (30, 2), (120,), (30,))

Import model from Library

```
In [77]: from sklearn.linear_model import LinearRegression # as slr  
mlr=LinearRegression()
```

Train Model with x & y train dataset

```
In [78]: mlr.fit(x_train,y_train)
```

Out[78]: LinearRegression()

Prediction with x_test and got predicted y

```
In [79]: y_pred=mlr.predict(x_test)  
y_pred[:5],y_test.values[:5]
```

```
Out[79]: (array([6.26455548, 4.74433166, 4.93077365, 5.43991048, 4.84902194]),  
 array([7. , 4.5, 5.4, 5.1, 4.8]))
```

Evaluation - Accuracy score

```
In [80]: mlr.score(x_test,y_test)
```

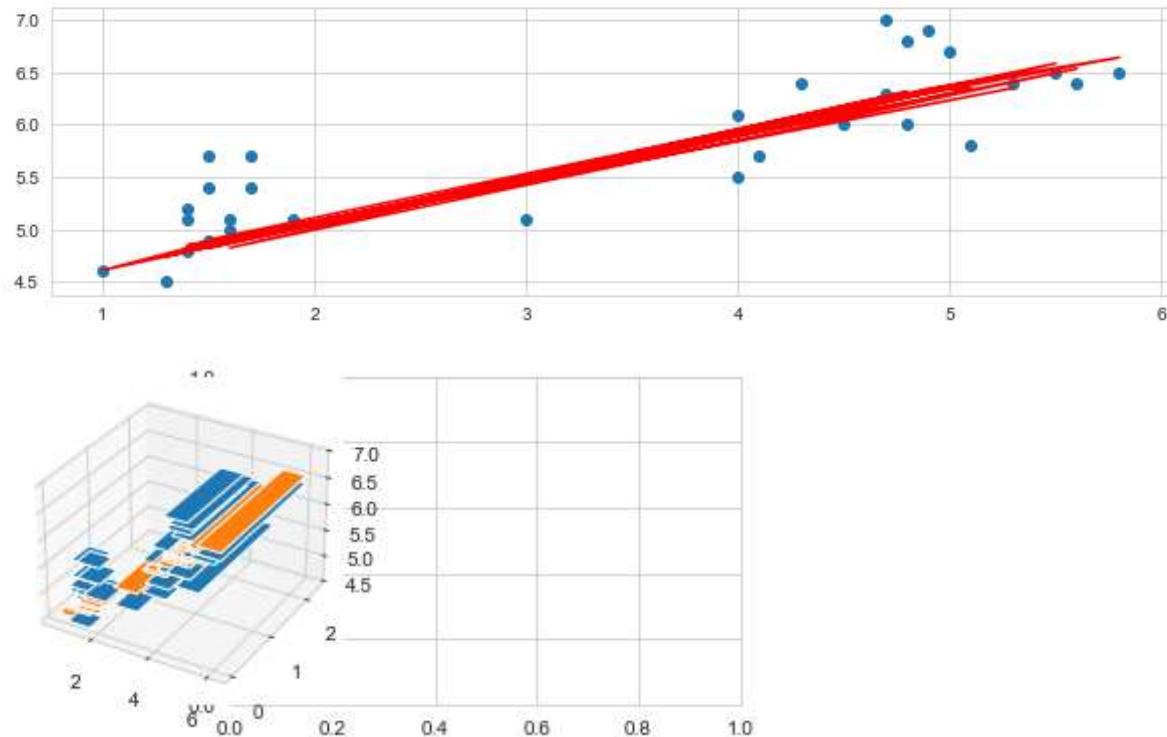
```
Out[80]: 0.7140431598884107
```

```
In [81]: from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score  
print(mean_absolute_error(y_test,y_pred))  
print(mean_squared_error(y_test,y_pred))  
print(r2_score(y_test,y_pred))
```

```
0.3028599497473194  
0.1450246001117043  
0.7140431598884107
```

Visualization

```
In [82]: plt.figure(figsize=(25,3))
plt.subplot(1,2,1)
plt.scatter(x_test['PetalLengthCm'],y_test)
plt.plot(x_test['PetalLengthCm'],y_pred,'red')
plt.show()
plt.figure(figsize=(10,3))
plt.subplot(1,2,2)
ax=plt.axes(projection='3d')
ax.bar(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.bar(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred)
plt.show()
ax.bar(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.bar(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred)
```



Out[82]: <BarContainer object of 30 artists>

In []:

Polynomial Regression -

```
In [83]: from sklearn.preprocessing import PolynomialFeatures
pf=PolynomialFeatures(degree=2)
x_train_pf=pf.fit_transform(x_train)
x_test_pf=pf.fit_transform(x_test)
```

```
In [84]: x_test.head(2)
```

Out[84]:

	PetalLengthCm	PetalWidthCm
50	4.7	1.4
41	1.3	0.3

<https://www.analyticsvidhya.com/blog/2021/07/all-you-need-to-know-about-polynomial-regression/#:~:text=Polynomial%20regression%20is%20a%20form%20of%20Linear%20regression>
<https://www.analyticsvidhya.com/blog/2021/07/all-you-need-to-know-about-polynomial-regression/#:~:text=Polynomial%20regression%20is%20a%20form%20of%20Linear%20regression>



```
In [85]: x_test_pf[:2]
```

```
Out[85]: array([[ 1. , 4.7 , 1.4 , 22.09, 6.58, 1.96],  
 [ 1. , 1.3 , 0.3 , 1.69, 0.39, 0.09]])
```

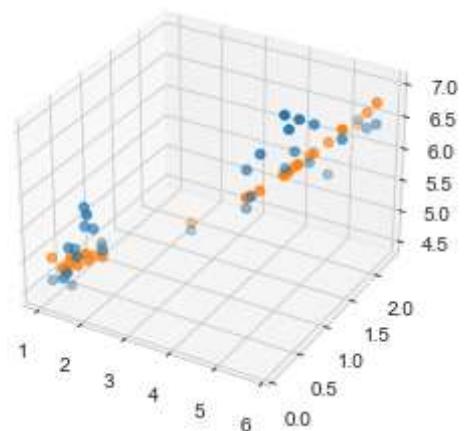
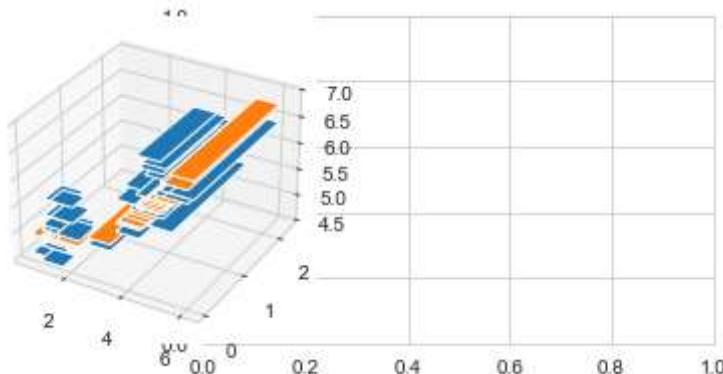
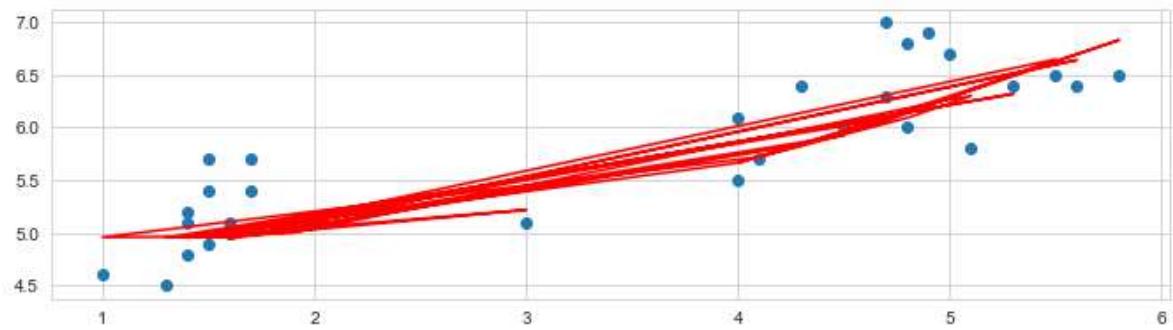
```
In [86]: from sklearn.linear_model import LinearRegression  
pr=LinearRegression()  
pr.fit(x_train_pf,y_train)  
y_pred_pr=pr.predict(x_test_pf)  
y_pred_pr[:5],y_test.values[:5]
```

```
Out[86]: (array([6.08103551, 4.95493418, 4.977728 , 5.21828609, 4.95334684]),  
 array([7. , 4.5, 5.4, 5.1, 4.8]))
```

```
In [87]: print('-----accuracy score-----')  
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score  
from math import sqrt  
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_pr))  
print('mean_squared_error',mean_squared_error(y_test,y_pred_pr))  
mse=mean_squared_error(y_test,y_pred_pr)  
print('r2_score',r2_score(y_test,y_pred_pr))  
print('MODEL SCORE',pr.score(x_test_pf,y_test))  
print(sqrt(mse))
```

```
-----accuracy score-----  
mean_absolute_error 0.3240276091521673  
mean_squared_error 0.16369077927756345  
r2_score 0.6772375310012113  
MODEL SCORE 0.6772375310012113  
0.4045871714199098
```

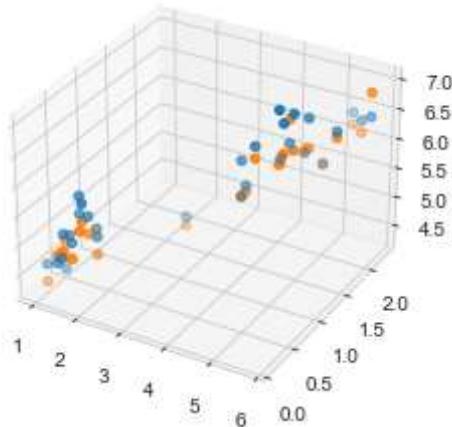
```
In [88]: plt.figure(figsize=(25,3))
plt.subplot(1,2,1)
plt.scatter(x_test['PetalLengthCm'],y_test)
plt.plot(x_test['PetalLengthCm'],y_pred_pr,'red')
plt.show()
plt.figure(figsize=(10,3))
plt.subplot(1,2,2)
ax=plt.axes(projection='3d')
ax.bar(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.bar(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_pr)
plt.show()
ax = plt.axes(projection ='3d')
fig = ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
fig = ax.scatter(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_pr,'red')
plt.show()
```



Decision tree Regressor

```
In [89]: from sklearn.tree import DecisionTreeRegressor
from sklearn.tree import DecisionTreeRegressor
dtr=DecisionTreeRegressor()
print(dtr.fit(x_train,y_train))
y_pred_dtr=dtr.predict(x_test)
print(y_pred_dtr[:5])
print(y_test.values[:5])
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_dtr))
print('mean_squared_error',mean_squared_error(y_test,y_pred_dtr))
mse=mean_squared_error(y_test,y_pred_dtr)
print('r2_score',r2_score(y_test,y_pred_dtr))
print('MODEL SCORE',dtr.score(x_test,y_test))
print(sqrt(mse))
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_dtr,'red')
plt.show()
```

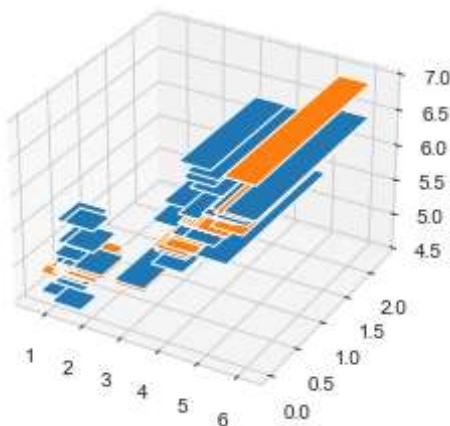
```
DecisionTreeRegressor()
[6.1      5.      5.1      4.95      4.92857143]
[7.  4.5  5.4  5.1  4.8]
-----accuracy score-----
mean_absolute_error 0.29183333333333356
mean_squared_error 0.13057430272108855
r2_score 0.7425359906034097
MODEL SCORE 0.7425359906034097
0.3613506644813159
```



RandomForestRegressor

```
In [90]: from sklearn.ensemble import RandomForestRegressor
rfr=RandomForestRegressor(n_estimators=5,max_depth=3)
rfr.fit(x_train,y_train)
y_pred_rfr=rfr.predict(x_test)
print(y_pred_rfr[:5],y_test.values[:5])
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_rfr))
print('mean_squared_error',mean_squared_error(y_test,y_pred_rfr))
mse=mean_squared_error(y_test,y_pred_rfr)
print('r2_score',r2_score(y_test,y_pred_rfr))
print('MODEL SCORE',rfr.score(x_test,y_test))
print(sqrt(mse))
ax = plt.axes(projection ='3d')
ax.bar(x_test[ 'PetalLengthCm'],x_test[ 'PetalWidthCm'],y_test)
ax.bar(x_test[ 'PetalLengthCm'],x_test[ 'PetalWidthCm'],y_pred_rfr)
plt.show()
```

```
[6.17281621 5.02102015 5.05012332 5.05012332 4.94959158] [7.  4.5 5.4 5.1 4.
8]
-----accuracy score-----
mean_absolute_error 0.3080258431088217
mean_squared_error 0.15257221813355867
r2_score 0.699160905441673
MODEL SCORE 0.699160905441673
0.39060493869581153
```

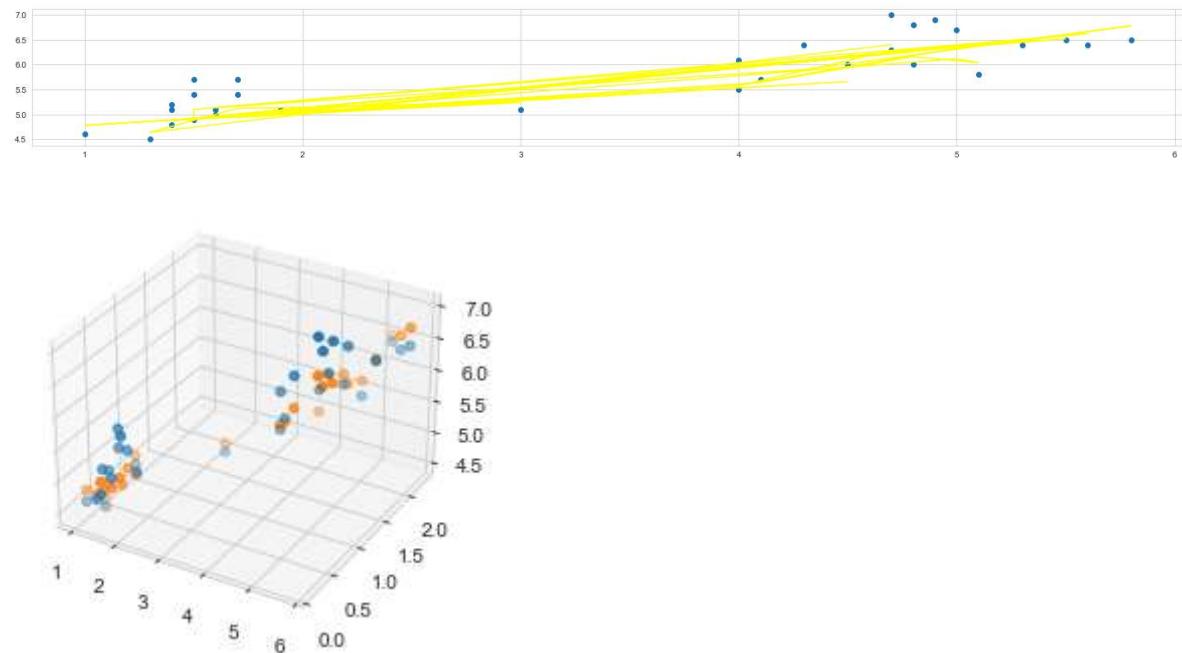


K-Nearest_neighbor

```
In [91]: from sklearn.neighbors import KNeighborsRegressor
knn=KNeighborsRegressor(n_neighbors=5)
knn.fit(x_train,y_train)
y_pred_knn=knn.predict(x_test)
print(y_pred_knn[:5],y_test.values[:5])
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_knn))
print('mean_squared_error',mean_squared_error(y_test,y_pred_knn))
print('r2_score',r2_score(y_test,y_pred_knn))
print('MODEL SCORE',knn.score(x_test,y_test))
```

[6.4 4.64 5.12 5.24 4.88] [7. 4.5 5.4 5.1 4.8]
-----accuracy score-----
mean_absolute_error 0.29399999999999993
mean_squared_error 0.13582666666666662
r2_score 0.7321794759442644
MODEL SCORE 0.7321794759442644

```
In [92]: plt.figure(figsize=(25,3))
plt.scatter(x_test['PetalLengthCm'],y_test)
plt.plot(x_test['PetalLengthCm'],y_pred_knn,'yellow')
plt.show()
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_knn)
plt.show()
```



Support Vector Machine - Regressor

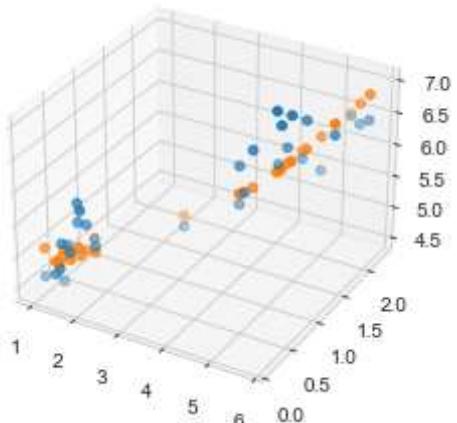
- https://www.saedsayad.com/support_vector_machine_reg.htm
[\(https://www.saedsayad.com/support_vector_machine_reg.htm\)](https://www.saedsayad.com/support_vector_machine_reg.htm)
- <https://www.analyticsvidhya.com/blog/2020/03/support-vector-regression-tutorial-for-machine-learning/>
[\(https://www.analyticsvidhya.com/blog/2020/03/support-vector-regression-tutorial-for-machine-learning/\)](https://www.analyticsvidhya.com/blog/2020/03/support-vector-regression-tutorial-for-machine-learning/)

```
In [93]: from sklearn.svm import SVR
```

```
In [94]: svr=SVR()
svr.fit(x_train,y_train)
y_pred_svr=svr.predict(x_test)
print(y_pred_svr[:5],y_test.values[:5])
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_svr))
print('mean_squared_error',mean_squared_error(y_test,y_pred_svr))
mse=mean_squared_error(y_test,y_pred_svr)
print('r2_score',r2_score(y_test,y_pred_svr))
print('MODEL SCORE',svr.score(x_test,y_test))
print(sqrt(mse))
```

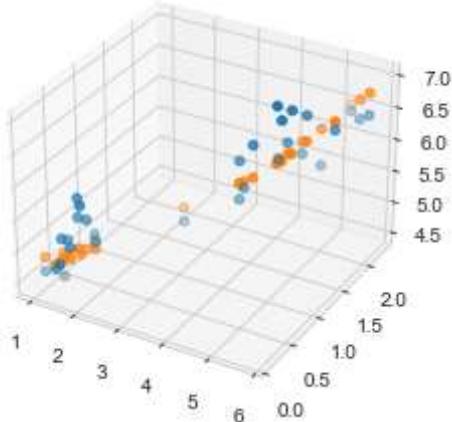
```
[6.06661625 5.01000585 4.99794693 5.27451893 5.00191259] [7. 4.5 5.4 5.1 4.8]
-----accuracy score-----
mean_absolute_error 0.3321669820345876
mean_squared_error 0.17025036123222761
r2_score 0.6643034679059574
MODEL SCORE 0.6643034679059574
0.41261405845199656
```

```
In [95]: ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_svr,'yellow')
plt.show()
```



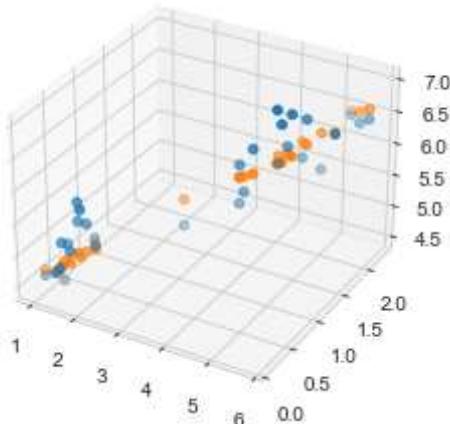
```
In [96]: print(' svc - poly')
svcp=SVR(kernel='poly', degree=2, C=.1)
svcp.fit(x_train,y_train)
y_pred_svcp=svcp.predict(x_test)
print(f'predicted_y-{y_pred_svcp[:5]}')
print(f'actual_y-{y_test.values[:5]}')
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_svcp))
print('mean_squared_error',mean_squared_error(y_test,y_pred_svcp))
mse=mean_squared_error(y_test,y_pred_svcp)
print('r2_score',r2_score(y_test,y_pred_svcp))
print('MODEL SCORE',svcp.score(x_test,y_test))
print(sqrt(mse))
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_svcp,'yellow')
plt.show()
```

```
svc - poly
predicted_y-[6.12104874 4.86709643 4.94047719 5.31962972 4.88163822]
actual_y-[7.  4.5 5.4 5.1 4.8]
-----accuracy score-----
mean_absolute_error 0.3299633765603468
mean_squared_error 0.16463171630025422
r2_score 0.6753822086797197
MODEL SCORE 0.6753822086797197
0.4057483410936565
```



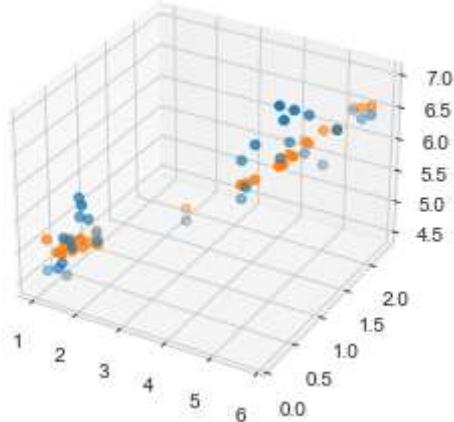
```
In [97]: print(' svc - linear')
svcl=SVR(kernel='linear', degree=2, C=.1).fit(x_train,y_train)
y_pred_svcl=svcl.predict(x_test)
print(f'predicted_y-[{y_pred_svcl[:5]}]')
print(f'actual_y-[{y_test.values[:5]}]')
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_svcl))
print('mean_squared_error',mean_squared_error(y_test,y_pred_svcl))
mse=mean_squared_error(y_test,y_pred_svcl)
print('r2_score',r2_score(y_test,y_pred_svcl))
print('MODEL SCORE',svcl.score(x_test,y_test))
print(sqrt(mse))
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_svcl,'red')
plt.show()
```

```
svc - linear
predicted_y-[6.18120077 4.80672237 4.96644849 5.51076795 4.83152816]
actual_y-[7.  4.5 5.4 5.1 4.8]
-----accuracy score-----
mean_absolute_error 0.31253644964643434
mean_squared_error 0.15319478502714945
r2_score 0.6979333394872613
MODEL SCORE 0.6979333394872613
0.39140105394230795
```



```
In [98]: print(' svc - rbf')
svcrbf=SVR(kernel='rbf', degree=2, C=.1)
svcrbf.fit(x_train,y_train)
y_pred_svcrbf=svcrbf.predict(x_test)
print(f'predicted_y-{y_pred_svcrbf[:5]}')
print(f'actual_y-{y_test.values[:5]}')
print('-----accuracy score-----')
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
from math import sqrt
print('mean_absolute_error',mean_absolute_error(y_test,y_pred_svcrbf))
print('mean_squared_error',mean_squared_error(y_test,y_pred_svcrbf))
mse=mean_squared_error(y_test,y_pred_svcrbf)
print('r2_score',r2_score(y_test,y_pred_svcrbf))
print('MODEL SCORE',svcl.score(x_test,y_test))
print(sqrt(mse))
ax = plt.axes(projection ='3d')
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_test)
ax.scatter3D(x_test['PetalLengthCm'],x_test['PetalWidthCm'],y_pred_svcrbf,'yellow')
plt.show()
```

```
svc - rbf
predicted_y-[6.07551106 5.05798193 5.03970157 5.29232494 5.05507296]
actual_y-[7.  4.5 5.4 5.1 4.8]
-----accuracy score-----
mean_absolute_error 0.3094455623876042
mean_squared_error 0.1580751891626363
r2_score 0.6883102483428871
MODEL SCORE 0.6979333394872613
0.39758670647122535
```



Unsupervised Learning

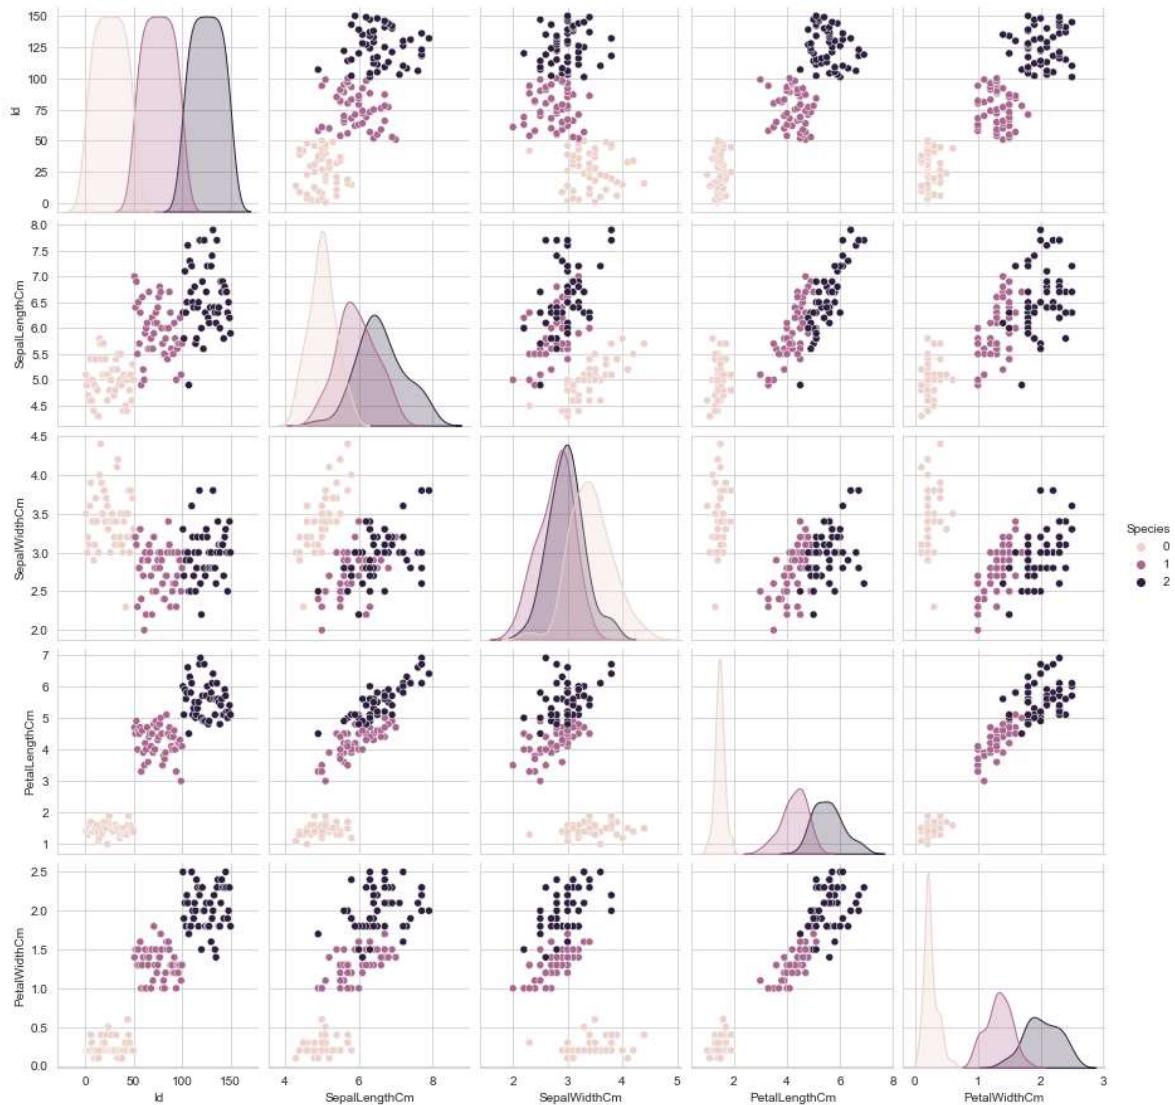
- K - Mean Cluster

Finding optimum no. of clusters

- K-Means clustering aims to partition data into k clusters in a way that data points in the same cluster are similar and data points in the different clusters are farther apart
- K-means clustering is unsupervised learning algorithm. Unsupervised learning algorithm is machine learning algorithm to analyze and cluster unlabeled datasets
- In this notebook we are going to cluster iris Dataset.

Data Set

```
In [99]: sns.pairplot(data=df1,hue='Species')
plt.show()
```



```
In [100]: df.head(2),df['Species'].value_counts() # with label encoding on species
```

```
Out[100]: (   Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
0    1           5.1          3.5          1.4          0.2      0
1    2           4.9          3.0          1.4          0.2      0,
0    50
1    50
2    50
Name: Species, dtype: int64)
```

```
In [101]: df1.head(2),df1['Species'].value_counts() # without label encoder
```

```
Out[101]: (   Id SepalLengthCm SepalWidthCm PetalLengthCm PetalWidthCm Species
0    1           5.1          3.5          1.4          0.2      0
1    2           4.9          3.0          1.4          0.2      0,
0    50
1    50
2    50
Name: Species, dtype: int64)
```

```
In [ ]:
```

Change data set into numpy arrays

```
In [102]: dfkm=df.values           # dataframe convert into numpy array
print(dfkm[:3])
```

```
[[1.  5.1 3.5 1.4 0.2 0. ]
 [2.  4.9 3.  1.4 0.2 0. ]
 [3.  4.7 3.2 1.3 0.2 0. ]]
```

Define features/independent/ x variable and target/dependent y variable

```
In [103]: x=dfkm[:,1:5]      # feature
y=dfkm[:, -1]       # target
x[:3],y[:3]
```

```
Out[103]: (array([[5.1, 3.5, 1.4, 0.2],
                   [4.9, 3., 1.4, 0.2],
                   [4.7, 3.2, 1.3, 0.2]]),
array([0., 0., 0.]))
```

Import K_Means Model / Algorithm

```
In [104]: from sklearn import cluster  
cluster.KMeans
```

```
Out[104]: sklearn.cluster._kmeans.KMeans
```

```
In [105]: cluster.KMeans(n_clusters=3)
```

```
Out[105]: KMeans(n_clusters=3)
```

```
km=cluster.KMeans(n_clusters=3,init='k-  
means++',max_iter=300,n_init=10,random_state=0)  
centroid=km.fit_predict(df.values[:,1:5])
```

Import K_Means Model / Algorithm

```
In [106]: from sklearn.cluster import KMeans
```

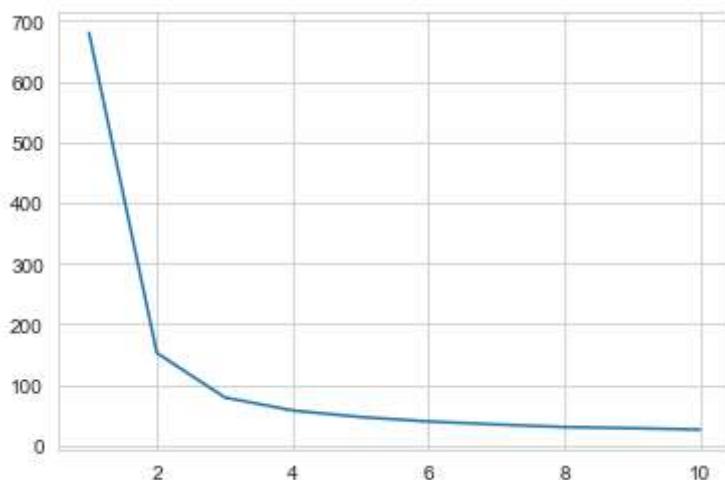
Defining cluster from elbow curve method

In [107]:

```
l=[]
for i in range(1,11):
    km=cluster.KMeans(i)
    km.fit(df.values[:,1:5])
    l.append(km.inertia_)
plt.plot(range(1,11),l[:10])
plt.show()
```

C:\Users\omkan\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:881: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.

```
warnings.warn(
```



In [108]:

```
km.inertia_
```

Out[108]:

```
25.909045084301678
```

<https://stackoverflow.com/questions/69596239/how-to-avoid-memory-leak-when-dealing-with-kmeans-for-example-in-this-code-i-am> (<https://stackoverflow.com/questions/69596239/how-to-avoid-memory-leak-when-dealing-with-kmeans-for-example-in-this-code-i-am>)

This is 'elbow curve method', it is widely used to identify optimum no. of clusters.

The basic idea is for each k value, we calculate average distances to the centroid across all data points.

The curve looks like an elbow. When average distances to the centroid doesn't decrease significantly with every iteration, we can choose optimum number of clusters. In the above plot, the elbow is at k=3.

Assign cluster

```
In [109]: k_m=KMeans(n_clusters=3)      # define 3 cluster/group/class here  
k_m
```

```
Out[109]: KMeans(n_clusters=3)
```

Fitting the x/feature dataset here

```
In [110]: km=k_m.fit(x)
```

```
In [111]: km
```

```
Out[111]: KMeans(n_clusters=3)
```

K-Mean algorithm clustering - output/target/y - Species

```
In [112]: km_predict=km.labels_ # predicted categories of species  
km_predict
```

```
Out[112]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
1, 1, 1, 1, 1, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
0, 0, 0, 2, 2, 0, 0, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 0, 0,  
0, 2, 0, 0, 0, 2, 0, 0, 2, 0, 0, 2, 0, 0, 2, 0, 0, 2, 0, 0, 2])
```

```
In [113]: center_point=km.cluster_centers_      # Center point of each features with target
```

```
In [114]: df1.head(2),df1['Species'].value_counts() # Associate with km.cluster_center
```

```
Out[114]: (   Id  SepalLengthCm  SepalWidthCm  PetalLengthCm  PetalWidthCm  Species  
0    1           5.1            3.5            1.4            0.2            0  
1    2           4.9            3.0            1.4            0.2            0,  
0    50  
1    50  
2    50  
Name: Species, dtype: int64)
```

```
In [115]: centroid=km_predict # K_mean predict y / target values  
centroid          # now centroid has - setosa,versicolor,virginica as 0,1,2
```

```
Out[115]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,  
1, 1, 1, 1, 1, 2, 2, 0, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2, 2,  
0, 0, 0, 2, 2, 0, 0, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 2, 0, 0, 0,  
0, 2, 0, 0, 0, 2, 0, 0, 2, 0, 0, 2, 0, 0, 2, 0, 0, 2, 0, 0, 2])
```

```
In [116]: x[:,1]      # all rows with sepallenthcm column
```

```
Out[116]: array([3.5, 3. , 3.2, 3.1, 3.6, 3.9, 3.4, 3.4, 2.9, 3.1, 3.7, 3.4, 3. ,  
3. , 4. , 4.4, 3.9, 3.5, 3.8, 3.8, 3.4, 3.7, 3.6, 3.3, 3.4, 3. ,  
3.4, 3.5, 3.4, 3.2, 3.1, 3.4, 4.1, 4.2, 3.1, 3.2, 3.5, 3.1, 3. ,  
3.4, 3.5, 2.3, 3.2, 3.5, 3.8, 3. , 3.8, 3.2, 3.7, 3.3, 3.2, 3.2,  
3.1, 2.3, 2.8, 2.8, 3.3, 2.4, 2.9, 2.7, 2. , 3. , 2.2, 2.9, 2.9,  
3.1, 3. , 2.7, 2.2, 2.5, 3.2, 2.8, 2.5, 2.8, 2.9, 3. , 2.8, 3. ,  
2.9, 2.6, 2.4, 2.4, 2.7, 2.7, 3. , 3.4, 3.1, 2.3, 3. , 2.5, 2.6,  
3. , 2.6, 2.3, 2.7, 3. , 2.9, 2.9, 2.5, 2.8, 3.3, 2.7, 3. , 2.9,  
3. , 3. , 2.5, 2.9, 2.5, 3.6, 3.2, 2.7, 3. , 2.5, 2.8, 3.2, 3. ,  
3.8, 2.6, 2.2, 3.2, 2.8, 2.8, 2.7, 3.3, 3.2, 2.8, 3. , 2.8, 3. ,  
2.8, 3.8, 2.8, 2.8, 2.6, 3. , 3.4, 3.1, 3. , 3.1, 3.1, 3.1, 2.7,  
3.2, 3.3, 3. , 2.5, 3. , 3.4, 3. ])
```

```
In [117]: x[centroid==0,1]      # x[ centroid == setosa,sepallength]
```

```
Out[117]: array([3.1, 3. , 3.3, 3. , 2.9, 3. , 3. , 2.9, 2.5, 3.6, 3.2, 2.7, 3. ,  
3.2, 3. , 3.8, 2.6, 3.2, 2.8, 3.3, 3.2, 2.8, 3. , 2.8, 3.8, 2.8,  
2.6, 3. , 3.4, 3.1, 3.1, 3.1, 3.2, 3.3, 3. , 3. , 3.4])
```

```
In [118]: km.cluster_centers_
```

```
Out[118]: array([[6.85      , 3.07368421, 5.74210526, 2.07105263],  
[5.006      , 3.418      , 1.464      , 0.244      ],  
[5.9016129 , 2.7483871 , 4.39354839, 1.43387097]])
```

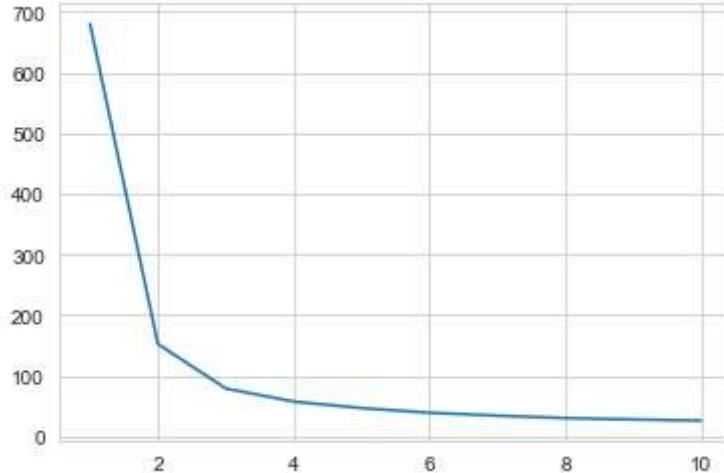
Plotting Prediction

```
In [119]: km.cluster_centers_[:,0],km.cluster_centers_[:,1]
```

```
Out[119]: (array([6.85      , 5.006      , 5.9016129]),  
array([3.07368421, 3.418      , 2.7483871 ]))
```

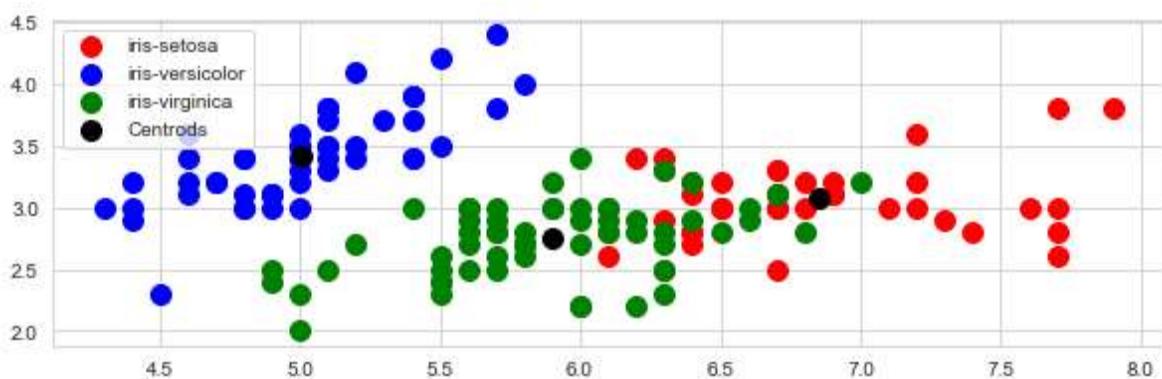
```
In [120]: centroid_list=[]
no_of_clusters=range(1,11)
for i in no_of_clusters:
    kmeans=cluster.KMeans(n_clusters=i,init='k-means++',max_iter=300,n_init=10)
    kmeans.fit(df.values[:,1:5])
    centroid_list.append(kmeans.inertia_)
plt.plot(no_of_clusters,centroid_list)
plt.show()
```

C:\Users\omkan\anaconda3\lib\site-packages\sklearn\cluster_kmeans.py:881: UserWarning: KMeans is known to have a memory leak on Windows with MKL, when there are less chunks than available threads. You can avoid it by setting the environment variable OMP_NUM_THREADS=1.
`warnings.warn(`



```
In [121]: plt.figure(figsize=(10,3))
plt.scatter(x[centroid==0,0],x[centroid==0,1],s=100,c='red',label='iris-setosa')
plt.scatter(x[centroid==1,0],x[centroid==1,1],s=100,c='blue',label='iris-versicolor')
plt.scatter(x[centroid==2,0],x[centroid==2,1],s=100,c='green',label='iris-virginica')

plt.scatter(km.cluster_centers_[:,0],km.cluster_centers_[:,1],s=100,c='black')
plt.legend()
plt.show()
```



```
In [122]: km.labels_
```

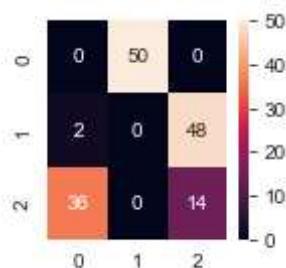
Checking our result

```
In [123]: import pandas as pd  
kmp=pd.crosstab(df1.values[:, -1], km.labels_)  
print(kmp)
```

<u>col_0</u>	0	1	2
<u>row_0</u>			
0.0	0	50	0
1.0	2	0	48
2.0	36	0	14

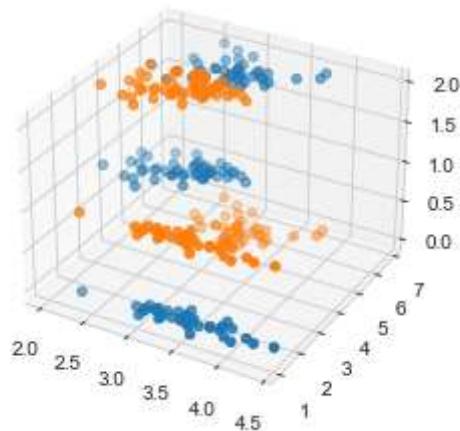
```
In [124]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_
cm=confusion_matrix(df.values[:, -1], km.labels_)
print(accuracy_score(df.values[:, -1], km.labels_))
print(classification_report(df.values[:, -1], km.labels_))
plt.figure(figsize=(2,2))
sns.heatmap(cm, annot=True)
plt.show()
```

	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	50
1.0	0.00	0.00	0.00	50
2.0	0.23	0.28	0.25	50
accuracy			0.09	150
macro avg	0.08	0.09	0.08	150
weighted avg	0.08	0.09	0.08	150



In []:

```
In [125]: ax = plt.axes(projection ='3d')
ax.scatter3D(x[:,1],x[:,2],y)
ax.scatter3D(x[:,1],x[:,2],km.labels_, 'yellow')
plt.show()
```

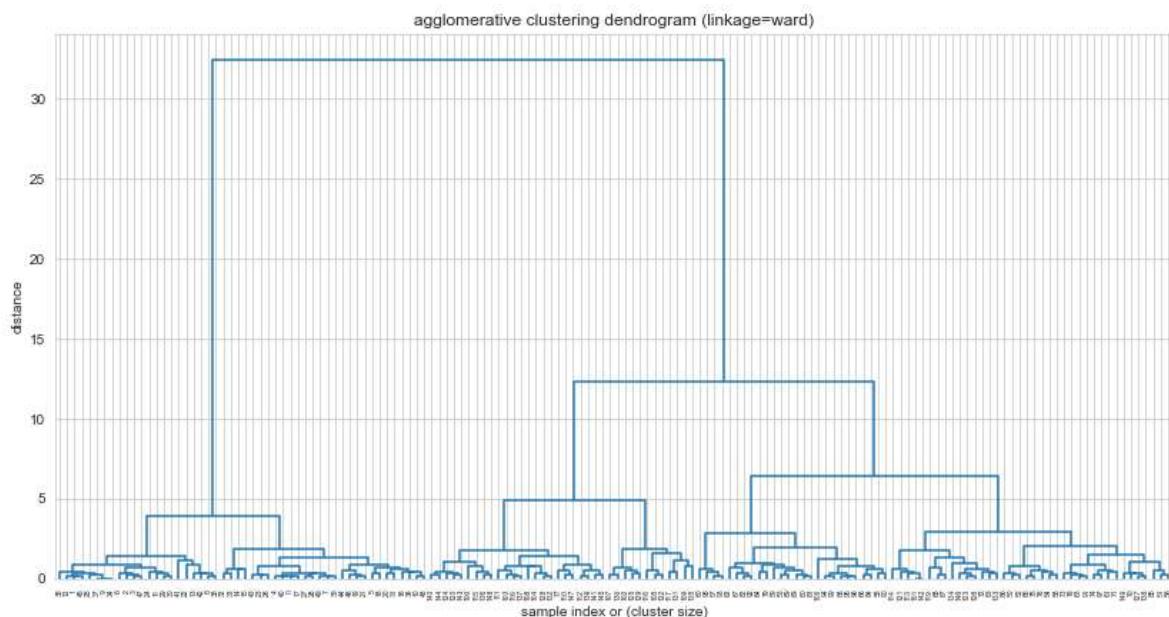


Hierachical Clustering

```
In [126]: from sklearn import datasets
from sklearn.cluster import AgglomerativeClustering
from scipy.cluster.hierarchy import dendrogram, linkage
import matplotlib.pyplot as plt
```

```
In [127]: dfhc=df.values
dfhc[:5]
x=dfhc[:,1:5]
y=dfhc[:, -1]
```

```
In [128]: linkage_matrix = linkage(x, 'ward')
plot = plt.figure(figsize=(14,7))
dendrogram(linkage_matrix,color_threshold=0)
plt.title('agglomerative clustering dendrogram (linkage=ward)')
plt.xlabel('sample index or (cluster size)')
plt.ylabel('distance')
plt.show()
plt.show()
```



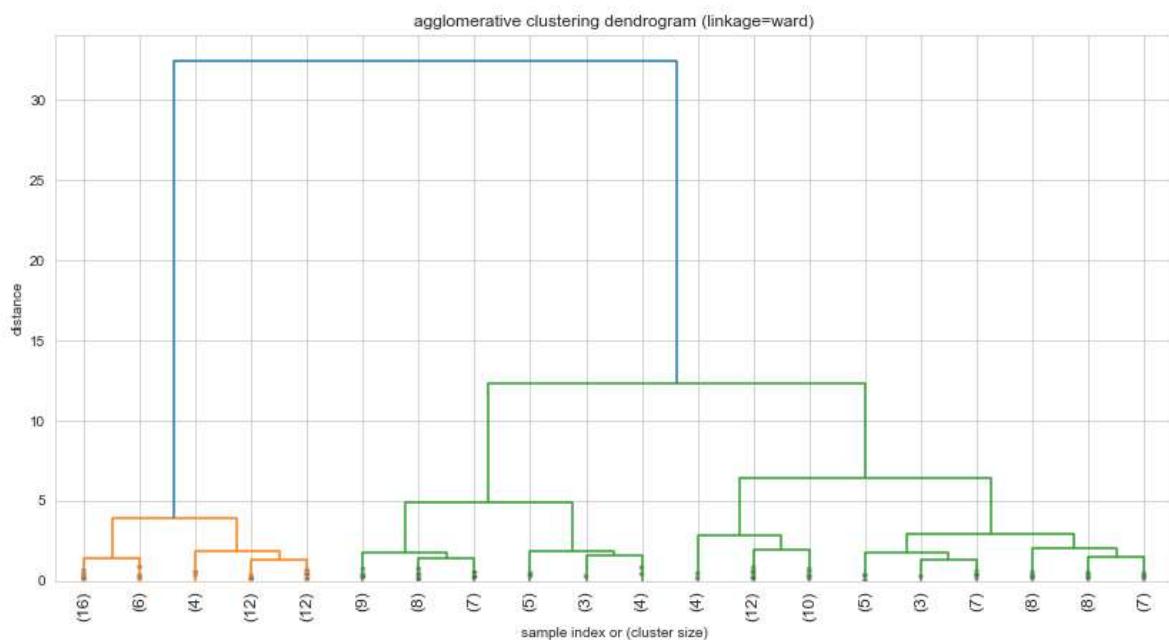
```
In [129]: hc=AgglomerativeClustering(linkage='ward', n_clusters=3)
```

```
In [130]: hc.fit(x)
hc_labels = hc.labels_
```

```
In [131]: hc_labels
```

```
Out[131]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
   1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
   1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
   0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
   0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
   2, 2, 2, 0, 0, 2, 2, 2, 2, 0, 2, 0, 2, 0, 2, 0, 2, 2, 2, 2, 0, 2, 2, 2,
   2, 0, 0, 2, 2, 2, 0, 2, 2, 2, 0, 2, 2, 0, 2, 2, 0, 2, 2, 0], dtype=int64)
```

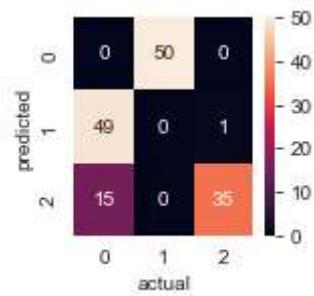
```
In [132]: plot_2 = plt.figure(figsize=(14,7))
dendrogram(linkage_matrix,truncate_mode='lastp',    #dendrogram(linkage_matrix
             p=20, # show only the last 20 merged clusters
             leaf_rotation=90., # rotates the x axis labels
             leaf_font_size=12.,# font size for x axis
             show_contracted=True) # to get a distribution impression
             # in truncated branches # color_threshold=0 )
plt.title('agglomerative clustering dendrogram (linkage=ward)')
plt.xlabel('sample index or (cluster size)')
plt.ylabel('distance')
plt.show()
```



```
In [133]: from sklearn.metrics import confusion_matrix, accuracy_score, classification_
cm=confusion_matrix(df.values[:, -1],hc.labels_)
print(accuracy_score(df.values[:, -1],hc.labels_))
print(classification_report(df.values[:, -1],hc.labels_))
plt.figure(figsize=(2,2))
sns.heatmap(cm, annot=True)
plt.xlabel('actual')
plt.ylabel('predicted')
plt.show()
import pandas as pd
kmp=pd.crosstab(df1.values[:, -1],hc.labels_)
print(kmp)
```

0.2333333333333334

	precision	recall	f1-score	support
0.0	0.00	0.00	0.00	50
1.0	0.00	0.00	0.00	50
2.0	0.97	0.70	0.81	50
accuracy			0.23	150
macro avg	0.32	0.23	0.27	150
weighted avg	0.32	0.23	0.27	150



col_0	0	1	2
row_0	0.0	0	50
1.0	49	0	1
2.0	15	0	35

Om Kant Sharma